Time Series (732A62) Lab1

Anubhav Dikshit(anudi287) and Maximilian Pfundstein(maxpf364)

09 September, 2019

Contents

Assignment 1. Computations with simulated data

a) Generate two time series $x_t = -0.8x_{t-2} + w_t$, where $x_0 = x_1 = 0$ and $x_t = cos(\frac{2\pi t}{5})$ with 100 observations each. Apply a smoothing filter $v_t = 0.2(x_t + x_{t-1} + x_{t-2} + x_{t-3} + x_{t-4})$ to these two series and compare how the filter has affected them.

```r
set.seed(12345)

n = 100
x <- vector(length = n)
x2 <- vector(length = n)

x[1] <- 0
x[2] <- 0

#first series generation
for(i in 3:n){
  x[i] <- -0.8 * x[i-2] + rnorm(1,0,1)
}

#second series generation
for(i in 1:n){
  x2[i] <- cos(0.4*pi*i)
}

# smoothing filter function
smoothing_filter <- function(x){
v <- vector(length = length(x))
for(i in 5:length(x)){
  v[i] = 0.2 * (x[i] + x[i-1] + x[i-2] + x[i-3] + x[i-4])
}
return(v)
}

#generate smoothed series
smooth_x <- smoothing_filter(x)
smooth_x2 <- smoothing_filter(x2)

#adding everything to a dataframe
df <- cbind(x,x2,smooth_x,smooth_x2) %>% as.data.frame() %>% mutate(index=1:100)

ggplot(df, aes(x=index)) +
  geom_line(aes(y=x, color="Original Time Series")) +
  geom_line(aes(y=smooth_x, color="Smoothened Time Series")) +
  ggtitle("Plot of 1st time series and its smoothened version") +
    scale_colour_manual("", breaks = c("Original Time Series", "Smoothened Time Series"),
                        values = c("#CC79A7", "#000000"))
```
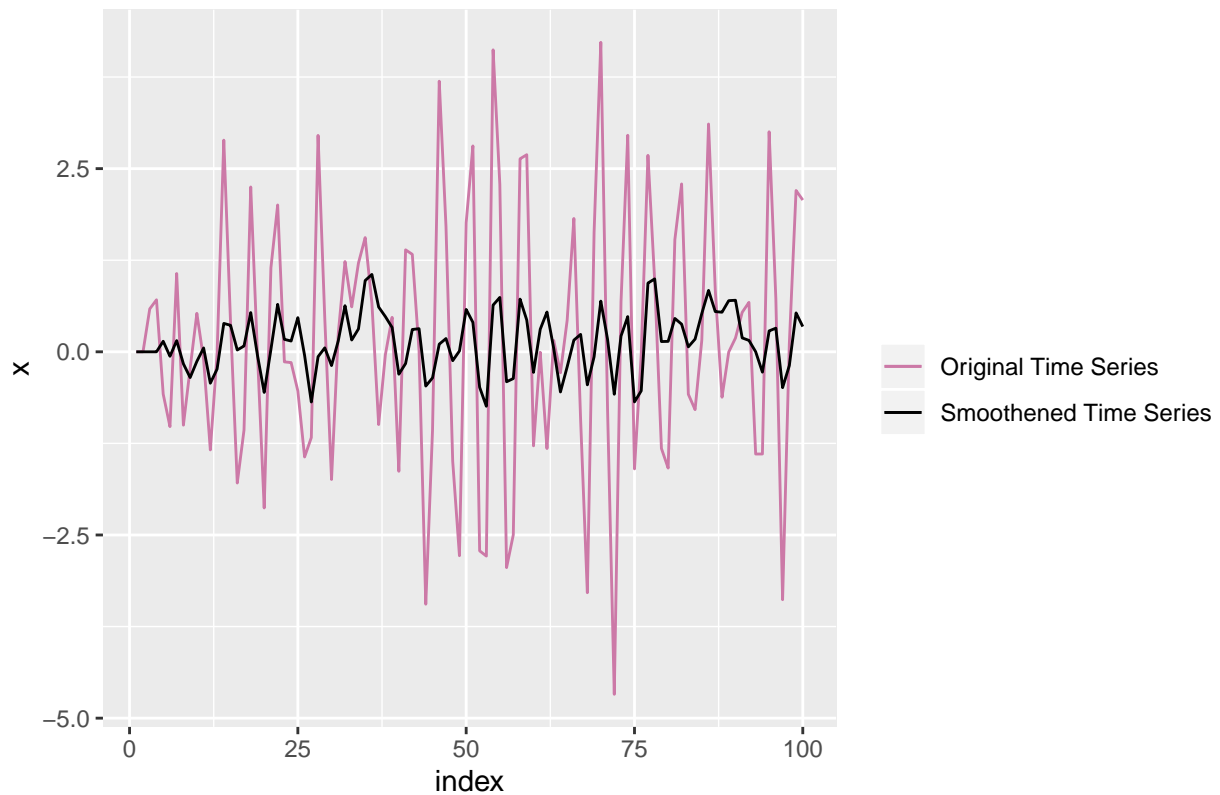
Plot of 1st time series and its smoothened version


```
ggplot(df, aes(x=index)) +
  geom_line(aes(y=x2, color="Original Time Series")) +
  geom_line(aes(y=smooth_x2, color="Smoothened Time Series")) +
  ggtitle("Plot of 2nd time series and its smoothened version") +
    scale_colour_manual("", breaks = c("Original Time Series", "Smoothened Time Series"),
                        values = c("#CC79A7", "#000000"))
```

Plot of 2nd time series and its smoothened version
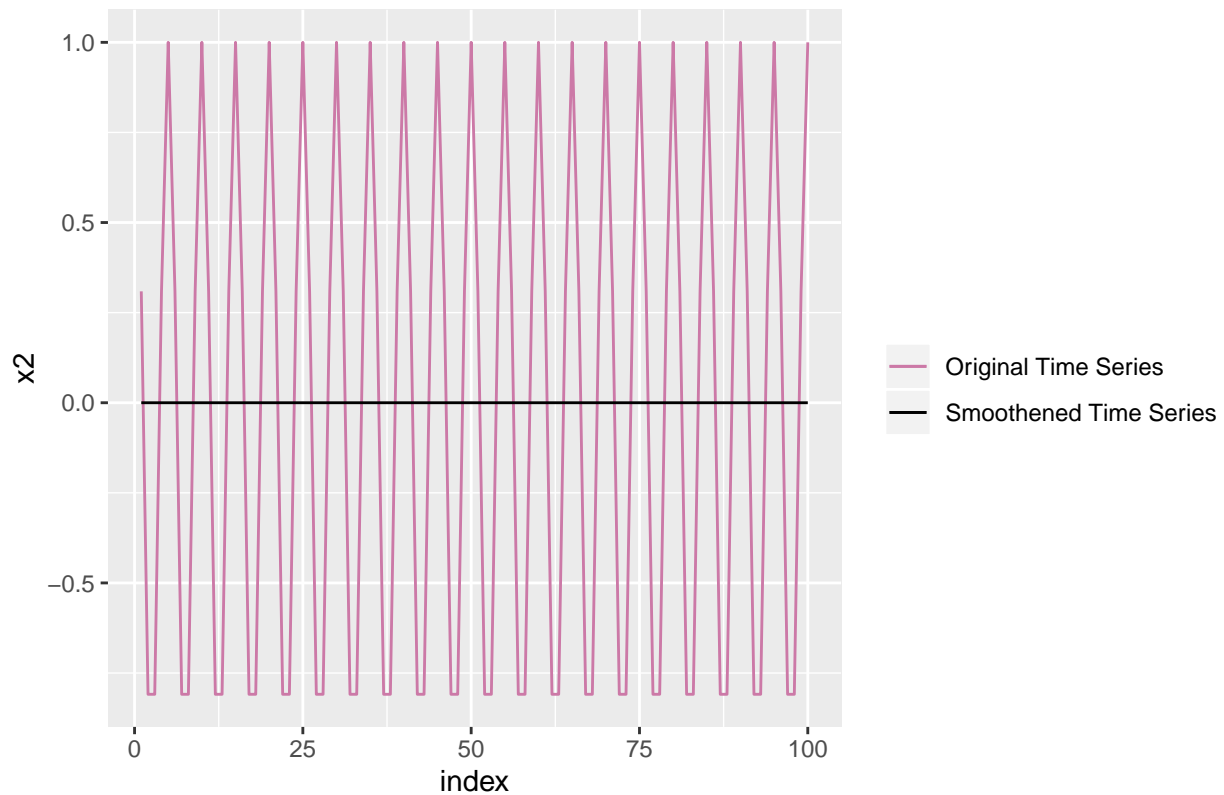
b) Consider time series $x_t - 4x_{t-1} + 2x_{t-2} + x_{t-5} = w_t + 3w_{t-2} + w_{t-4} - 4w_{t-6}$. Write an appropriate R code to investigate whether this time series is casual and invertible.

Causality: ARMA(p,q) is causal iff roots $\phi(z') = 0$ are outside unit circle. eg: $x_t = 0.4x_{t-1} + 0.3x_{t-2} + w_t$, roots are -> $1 - 0.4B + 0.3B^2$

equation is: $\phi(Z) = 1 - 4B + 2B^2 + 0B^3 + 0B^4 + B^5$

```
z = c(1,-4,2,0,0,1)
polyroot(z)
```

```
## [1]  0.2936658+0.000000i -1.6793817+0.000000i  1.0000000-0.000000i
## [4]  0.1928579-1.410842i  0.1928579+1.410842i
```

```
any(Mod(polyroot(z))<=1)
```

```
## [1] TRUE
```

Invertible: ARMA(p,q) is causal iff roots $\theta(z') = 0$ are outside unit circle.

equation is: $\theta(Z) = 1 + 3B^2 + B^4 - 4B^6$

```
z = c(1,0,3,0,1,0,-4)
polyroot(z)
```

```
## [1]  0.1375513+0.6735351i -0.1375513+0.6735351i -0.1375513-0.6735351i
## [4]  0.1375513-0.6735351i  1.0580446+0.0000000i -1.0580446+0.0000000i
```

```
any(Mod(polyroot(z))<=1)
```

[1] TRUE

Analysis: Baring one of the roots all are inside the unit circle. Thus the time series is not invertiable.

c) Use built-in R functions to simulate 100 observations from the process $x_t + \frac{3}{4}x_{t-1} = w_t - \frac{1}{9}w_{t-2}$ compute sample ACF and theoretical ACF, use seed 54321. Compare the ACF plots.

```
set.seed(54321)

series <- arima.sim(n = 100, list(ar = c(-3/4), ma = c(0,-1/9)))

acf(series)
```

Series series


```
acf(ARMAacf(ar = c(-3/4), ma =  c(0,-1/9), lag.max = 20))
```

5

Series ARMAacf(ar = c(−3/4), ma = c(0, −1/9), lag.max = 20)

Analysis: In the theortical ACF, only the 1 and 2nd lag components were significant, while using the sample ACF function we get many more lag components as significant.

Assignment 2. Visualization, detrending and residual analysis of Rhine data.

The data set Rhine.csv contains monthly concentrations of total nitrogen in the Rhine River in the period 1989-2002.

a) Import the data to R, convert it appropriately to ts object (use function ts()) and explore it by plotting the time series, creating scatter plots of x_t against $x_{t-1}, ...x_{t-12}$. Analyze the time series plot and the scatter plots: Are there any trends, linear or seasonal, in the time series? When during the year is the concentration highest? Are there any special patterns in the data or scatterplots? Does the variance seem to change over time? Which variables in the scatterplots seem to have a significant relation to each other?

```r
set.seed(12345)
rhine_data <- read.csv2("Rhine.csv")
rhine_data_ts <- ts(data = rhine_data$TotN_conc,
                     start = c(1989,1),
                     frequency = 12)

plot.ts(rhine_data_ts)
```

```
lag.plot(rhine_data_ts,lags = 12)
```
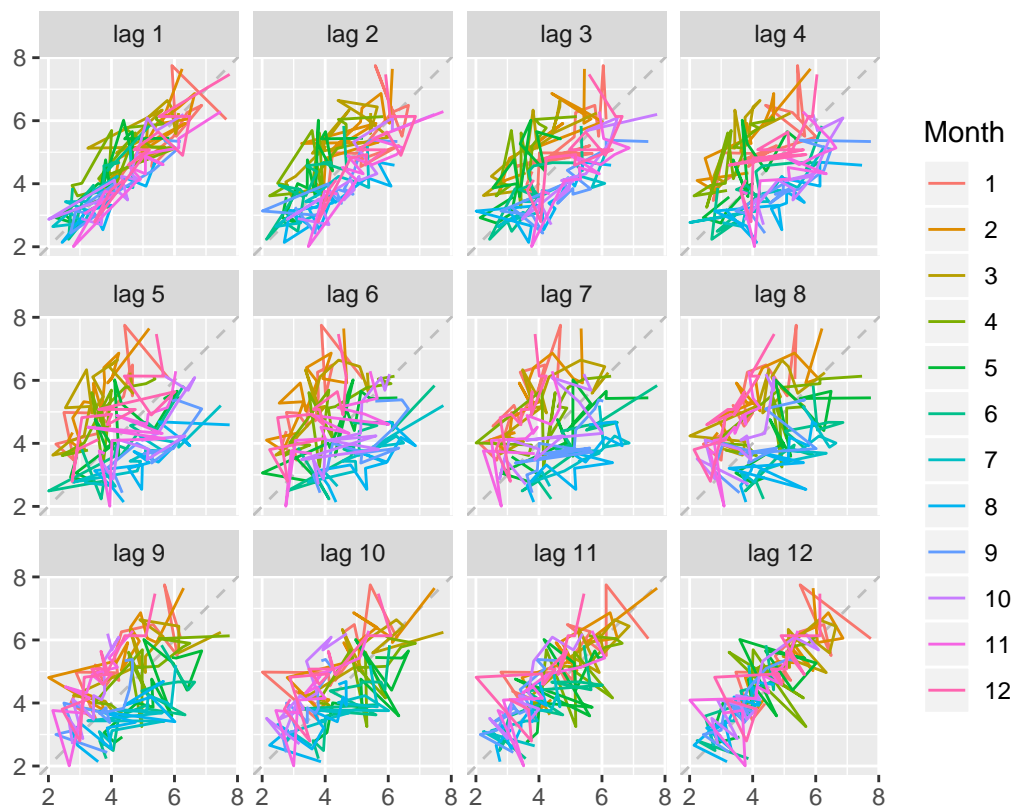
```
acf(rhine_data_ts)
```

Series rhine_data_ts


```
#alternative
autoplot(rhine_data_ts) + ylab("Total Concentration") +xlab("Year") +
  ggtitle("Concentration in Rhine vs. Year")
```

Concentration in Rhine vs. Year

```
gglagplot(rhine_data_ts, lags = 1, set.lags = 1:12, color=FALSE)
```

```
ggAcf(rhine_data_ts)
```

Series: rhine_data_ts

Analysis:

b) Eliminate the trend by fitting a linear model with respect to t to the time series. Is there a significant time trend? Look at the residual pattern and the sample ACF of the residuals and comment how this pattern might be related to seasonality of the series.

```
set.seed(12345)

rhine_lm_model <- lm(TotN_conc~Time, data=rhine_data)
plot(rhine_lm_model$residuals, type = 'l')
```

```
acf(rhine_lm_model$residuals)
```

Series rhine_lm_model$residuals

c) Eliminate the trend by fitting a kernel smoother with respect to t to the time series (choose a reasonable bandwidth yourself so the fit looks reasonable). Analyze the residual pattern and the sample ACF of the residuals and compare it to the ACF from step b). Conclusions? Do residuals seem to represent a stationary series?

```r
set.seed(12345)

model_smooth_lag_5 <- ksmooth(x = rhine_data$Time, y = rhine_data$TotN_conc, bandwidth=5)
model_smooth_lag_10 <- ksmooth(x = rhine_data$Time, y = rhine_data$TotN_conc, bandwidth=10)
model_smooth_lag_20 <- ksmooth(x = rhine_data$Time, y = rhine_data$TotN_conc, bandwidth=20)

model_smooth_lag_5_residual <- rhine_data$TotN_conc - model_smooth_lag_5$y
model_smooth_lag_10_residual <- rhine_data$TotN_conc - model_smooth_lag_10$y
model_smooth_lag_20_residual <- rhine_data$TotN_conc - model_smooth_lag_20$y

residual_df <- cbind(model_smooth_lag_5_residual, model_smooth_lag_10_residual,
                model_smooth_lag_20_residual, rhine_data$Time) %>% as.data.frame()

colnames(residual_df) <- c("lag_5_residual", "lag_10_residual", "lag_20_residual", "Time")

ggplot(residual_df, aes(x=Time)) +
  geom_line(aes(y=lag_5_residual, color="Lag 5 residual")) +
  geom_line(aes(y=lag_10_residual, color="Lag 10 residual")) +
  geom_line(aes(y=lag_20_residual, color="Lag 20 residual")) +
  ggtitle("Residual vs. Time by Lag") +
    scale_colour_manual("", breaks = c("Lag 5 residual", "Lag 10 residual", "Lag 20 residual"),
```

```
                    values = c("#CC79A7", "#000000", "#D55E00"))
```

Residual vs. Time by Lag


```
acf(model_smooth_lag_5_residual)
```

Series model_smooth_lag_5_residual


```
acf(model_smooth_lag_10_residual)
```

Series model_smooth_lag_10_residual


```r
acf(model_smooth_lag_20_residual)
```

Series model_smooth_lag_20_residual

---

d) Eliminate the trend by fitting the following so-called seasonal means model: $x_t = \alpha_0 + \alpha_1 t + \beta_1 I(month = 2) + ... + \beta_{12} I(month = 12) + w_t$, where I(x)=1 is an identity function. Fitting of this model will require you to augment data with a categorical variable showing the current month, and then fitting a usual linear regression. Analyze the residual pattern and the ACF of residuals.

```r
set.seed(12345)

rhine_data_wide <- rhine_data
rhine_data_wide$dummy <- "1"
rhine_data_wide$Month <- paste0("Month_",rhine_data_wide$Month)
rhine_data_wide <- dcast(rhine_data_wide,
                         formula = TotN_conc+Year+Time~Month, value.var = "dummy", fill = "0")

lm_model_month_lag <- lm(data=rhine_data_wide,
                    TotN_conc~Time+Month_1+Month_2+Month_3+Month_4+Month_5+Month_6+Month_7+
                      Month_8+Month_9+Month_10+Month_11+Month_12)


plot(lm_model_month_lag$residuals, type = 'l')
```

```
acf(lm_model_month_lag$residuals)
```

Series lm_model_month_lag$residuals

Analysis:

e) Perform stepwise variable selection in model from step d). Which model gives you the lowest AIC value? Which variables are left in the model?

```
set.seed(12345)

lm_model_month_lag_step <- stepAIC(lm_model_month_lag,
                              scope = list(upper = ~Time+Month_1+Month_2+
                                          Month_3+Month_4+Month_5+Month_6+Month_7+
                                          Month_8+Month_9+Month_10+Month_11+Month_12,
                                              lower = ~1), trace = TRUE,
                          direction="backward")
```

```
## Start:  AIC=-202.02
## TotN_conc ~ Time + Month_1 + Month_2 + Month_3 + Month_4 + Month_5 +
##     Month_6 + Month_7 + Month_8 + Month_9 + Month_10 + Month_11 +
##     Month_12
##
##
## Step:  AIC=-202.02
## TotN_conc ~ Time + Month_1 + Month_2 + Month_3 + Month_4 + Month_5 +
##     Month_6 + Month_7 + Month_8 + Month_9 + Month_10 + Month_11
##
##           Df Sum of Sq     RSS      AIC
## - Month_4  1     0.200  43.436 -203.249
```

```
## - Month_1    1      0.220  43.456 -203.170
## - Month_3    1      0.331  43.567 -202.743
## <none>                     43.237 -202.023
## - Month_2    1      1.440  44.677 -198.517
## - Month_11   1      2.305  45.541 -195.297
## - Month_5    1      3.274  46.511 -191.760
## - Month_10   1      3.401  46.637 -191.303
## - Month_9    1      7.853  51.089 -175.986
## - Month_6    1      8.215  51.452 -174.797
## - Month_7    1     14.321  57.557 -155.959
## - Month_8    1     16.488  59.725 -149.749
## - Time       1    118.387 161.624   17.499
##
## Step:  AIC=-203.25
## TotN_conc ~ Time + Month_1 + Month_2 + Month_3 + Month_5 + Month_6 +
##      Month_7 + Month_8 + Month_9 + Month_10 + Month_11
##
##             Df Sum of Sq     RSS      AIC
## <none>                     43.436 -203.249
## - Month_1    1      0.640  44.077 -202.790
## - Month_3    1      0.851  44.288 -201.988
## - Month_11   1      2.235  45.671 -196.819
## - Month_2    1      2.706  46.142 -195.096
## - Month_5    1      3.355  46.791 -192.748
## - Month_10   1      3.502  46.938 -192.223
## - Month_9    1      8.868  52.304 -174.036
## - Month_6    1      9.317  52.753 -172.602
## - Month_7    1     16.912  60.348 -150.004
## - Month_8    1     19.636  63.072 -142.586
## - Time       1    118.194 161.630   15.506
```

```r
colnames(lm_model_month_lag_step$model)
```

```
##  [1] "TotN_conc" "Time"      "Month_1"   "Month_2"   "Month_3"
##  [6] "Month_5"   "Month_6"   "Month_7"   "Month_8"   "Month_9"
## [11] "Month_10"  "Month_11"
```

Analysis: The final terms in the model are given above, this model had the least AIC.

Assignment 3. Analysis of oil and gas time series.

Weekly time series oil and gas present in the package astsa show the oil prices in dollars per barrel and gas prices in cents per dollar.

a) Plot the given time series in the same graph. Do they look like stationary series? Do the processes seem to be related to each other? Motivate your answer.

```r
set.seed(12345)

data_oil <- astsa::oil
data_gas <- astsa::gas

ts.plot(data_oil, data_gas, gpars = list(col = c("black", "red")))
```

```
#alternative

autoplot(ts(cbind(data_oil, data_gas), start = 2000, frequency = 52)) +
        ylab("Price of Oil and Gas") +xlab("Year") +
        ggtitle("Price of Oil and Gas vs. Years")
```

Price of Oil and Gas vs. Years

b) Apply log-transform to the time series and plot the transformed data. In what respect did this transformation made the data easier for the analysis?

```
set.seed(12345)

autoplot(ts(cbind(log(data_oil), log(data_gas)), start = 2000, frequency = 52)) +
        ylab("Log of price of Oil and Gas") +xlab("Year") +
        ggtitle("Log of price of Oil and Gas vs. Years")
```

Log of price of Oil and Gas vs. Years

c) To eliminate trend, compute the first difference of the transformed data, plot the detrended series, check their ACFs and analyze the obtained plots. Denote the data obtained here as x_t(oil) and y_t(gas).

```r
set.seed(12345)

autoplot(ts(diff(data_oil, differences = 1), start = 2000, frequency = 52)) +
        ylab("Price of Oil") +xlab("Year") +
        ggtitle("Price of Oil with Diff 1 vs. Years")
```

Price of Oil with Diff 1 vs. Years


```
autoplot(ts(diff(data_gas, differences = 1), start = 2000, frequency = 52)) +
        ylab("Price of Gas") +xlab("Year") +
        ggtitle("Price of Gas with Diff 1 vs. Years")
```

## Price of Gas with Diff 1 vs. Years


```
ggAcf(diff(data_oil, differences = 1), data_oil)
```

Series: diff(data_oil, differences = 1)

```
ggAcf(diff(data_gas, differences = 1), data_gas)
```

Series: diff(data_gas, differences = 1)

> d) Exhibit scatterplots of x_t and y_t for up to three weeks of lead time of x_t include a nonparametric smoother in each plot and comment the results: are there outliers? Are the relationships linear? Are there changes in the trend?

```r
set.seed(12345)

oil_price_one_diff <- diff(data_oil, differences = 1)
gas_price_one_diff <- diff(data_gas, differences = 1)

df <- data.frame(oil_price_one_diff=as.matrix(oil_price_one_diff),
        gas_price_one_diff = as.matrix(gas_price_one_diff),
                time=time(oil_price_one_diff))

df <- na.omit(df)


df$gas_price_one_diff = lag(df$gas_price_one_diff,1)
df$gas_price_two_diff = lag(df$gas_price_one_diff,2)
df$gas_price_three_diff = lag(df$gas_price_one_diff,3)
df$oil_price_one_diff = lag(df$oil_price_one_diff,1)
df$oil_price_two_diff = lag(df$oil_price_one_diff,2)
df$oil_price_three_diff = lag(df$oil_price_one_diff,3)

df <- na.omit(df)

df$smooth_one_week_lag <- ksmooth(x = df$oil_price_one_diff, y = df$gas_price_one_diff, bandwidth = 0.4
```

```
df$smooth_two_week_lag <- ksmooth(x = df$oil_price_two_diff, y = df$gas_price_two_diff, bandwidth = 0.4
df$smooth_three_week_lag <- ksmooth(x = df$oil_price_three_diff, y = df$gas_price_three_diff, bandwidth


ggplot(data=df, aes(x=time)) +
  geom_line(aes(y= smooth_one_week_lag, color= "smooth_one_week_lag")) +
    geom_line(aes(y= smooth_two_week_lag, color= "smooth_two_week_lag")) +
    geom_line(aes(y= smooth_three_week_lag, color= "smooth_three_week_lag")) +
      scale_colour_manual("", breaks = c("smooth_one_week_lag", "smooth_two_week_lag", "smooth_three_we
                          values = c("#CC79A7", "#000000", "#D55E00")) +
  ggtitle("Smoothened Plot of lag")
```


Smoothened Plot of lag

e) Fit the following model: $y_t = \alpha_0 + \alpha_1 I(x_t > 0) + \beta_1 x_1 + \beta_2 x_{t-1} + w_t$ and check which coefficients seem to be significant. How can this be interpreted? Analyze the residual pattern and the ACF of the residuals.

```
set.seed(12345)

df$x_t_more_zero <- ifelse(df$oil_price_one_diff>0,"1","0")
lm_model_lag <- lm(data=df, formula = gas_price_one_diff~oil_price_one_diff+oil_price_two_diff)

plot(lm_model_lag$residuals, type = 'l')
```

```
ggAcf(lm_model_lag$residuals)
```

Series: lm_model_lag$residuals

Appendix

```
knitr::opts_chunk$set(echo = TRUE)
options(scipen=999)

library("tidyverse") #ggplot and dplyr
library("gridExtra") # combine plots
library("knitr") # for pdf
library("fpp2") #timeseries with autoplot and stuff
library("reshape2") #reshape the data
library("MASS") #StepAIC
library("astsa") #dataset oil and gas is present here
library("zoo") #dataset oil and gas is present here


# The palette with black:
cbbPalette <- c("#000000", "#E69F00", "#56B4E9", "#009E73",
                "#F0E442", "#0072B2", "#D55E00", "#CC79A7")
set.seed(12345)
set.seed(12345)

n = 100
x <- vector(length = n)
```

```r
x2 <- vector(length = n)

x[1] <- 0
x[2] <- 0

#first series generation
for(i in 3:n){
  x[i] <- -0.8 * x[i-2] + rnorm(1,0,1)
}

#second series generation
for(i in 1:n){
  x2[i] <- cos(0.4*pi*i)
}

# smoothing filter function
smoothing_filter <- function(x){
v <- vector(length = length(x))
for(i in 5:length(x)){
  v[i] = 0.2 * (x[i] + x[i-1] + x[i-2] + x[i-3] + x[i-4])
}
return(v)
}

#generate smoothed series
smooth_x <- smoothing_filter(x)
smooth_x2 <- smoothing_filter(x2)

#adding everything to a dataframe
df <- cbind(x,x2,smooth_x,smooth_x2) %>% as.data.frame() %>% mutate(index=1:100)

ggplot(df, aes(x=index)) +
  geom_line(aes(y=x, color="Original Time Series")) +
  geom_line(aes(y=smooth_x, color="Smoothened Time Series")) +
  ggtitle("Plot of 1st time series and its smoothened version") +
    scale_colour_manual("", breaks = c("Original Time Series", "Smoothened Time Series"),
                        values = c("#CC79A7", "#000000"))

ggplot(df, aes(x=index)) +
  geom_line(aes(y=x2, color="Original Time Series")) +
  geom_line(aes(y=smooth_x2, color="Smoothened Time Series")) +
  ggtitle("Plot of 2nd time series and its smoothened version") +
    scale_colour_manual("", breaks = c("Original Time Series", "Smoothened Time Series"),
                        values = c("#CC79A7", "#000000"))

z = c(1,-4,2,0,0,1)
polyroot(z)
any(Mod(polyroot(z))<=1)
z = c(1,0,3,0,1,0,-4)
polyroot(z)
any(Mod(polyroot(z))<=1)
set.seed(54321)
```

```r
series <- arima.sim(n = 100, list(ar = c(-3/4), ma = c(0,-1/9)))

acf(series)
acf(ARMAacf(ar = c(-3/4), ma =  c(0,-1/9), lag.max = 20))
set.seed(12345)
rhine_data <- read.csv2("Rhine.csv")
rhine_data_ts <- ts(data = rhine_data$TotN_conc,
                     start = c(1989,1),
                     frequency = 12)

plot.ts(rhine_data_ts)
lag.plot(rhine_data_ts,lags = 12)
acf(rhine_data_ts)

#alternative
autoplot(rhine_data_ts) + ylab("Total Concentration") +xlab("Year") +
  ggtitle("Concentration in Rhine vs. Year")

gglagplot(rhine_data_ts, lags = 1, set.lags = 1:12, color=FALSE)
ggAcf(rhine_data_ts)

set.seed(12345)

rhine_lm_model <- lm(TotN_conc~Time, data=rhine_data)
plot(rhine_lm_model$residuals, type = 'l')
acf(rhine_lm_model$residuals)

set.seed(12345)

model_smooth_lag_5 <- ksmooth(x = rhine_data$Time, y = rhine_data$TotN_conc, bandwidth=5)
model_smooth_lag_10 <- ksmooth(x = rhine_data$Time, y = rhine_data$TotN_conc, bandwidth=10)
model_smooth_lag_20 <- ksmooth(x = rhine_data$Time, y = rhine_data$TotN_conc, bandwidth=20)

model_smooth_lag_5_residual <- rhine_data$TotN_conc - model_smooth_lag_5$y
model_smooth_lag_10_residual <- rhine_data$TotN_conc - model_smooth_lag_10$y
model_smooth_lag_20_residual <- rhine_data$TotN_conc - model_smooth_lag_20$y

residual_df <- cbind(model_smooth_lag_5_residual, model_smooth_lag_10_residual,
                     model_smooth_lag_20_residual, rhine_data$Time) %>% as.data.frame()

colnames(residual_df) <- c("lag_5_residual", "lag_10_residual", "lag_20_residual", "Time")

ggplot(residual_df, aes(x=Time)) +
  geom_line(aes(y=lag_5_residual, color="Lag 5 residual")) +
  geom_line(aes(y=lag_10_residual, color="Lag 10 residual")) +
  geom_line(aes(y=lag_20_residual, color="Lag 20 residual")) +
  ggtitle("Residual vs. Time by Lag") +
    scale_colour_manual("", breaks = c("Lag 5 residual", "Lag 10 residual", "Lag 20 residual"),
                        values = c("#CC79A7", "#000000", "#D55E00"))


acf(model_smooth_lag_5_residual)
acf(model_smooth_lag_10_residual)
```

```r
acf(model_smooth_lag_20_residual)
set.seed(12345)

rhine_data_wide <- rhine_data
rhine_data_wide$dummy <- "1"
rhine_data_wide$Month <- paste0("Month_",rhine_data_wide$Month)
rhine_data_wide <- dcast(rhine_data_wide,
                         formula = TotN_conc+Year+Time~Month, value.var = "dummy", fill = "0")

lm_model_month_lag <- lm(data=rhine_data_wide,
                    TotN_conc~Time+Month_1+Month_2+Month_3+Month_4+Month_5+Month_6+Month_7+
                    Month_8+Month_9+Month_10+Month_11+Month_12)


plot(lm_model_month_lag$residuals, type = 'l')
acf(lm_model_month_lag$residuals)

set.seed(12345)

lm_model_month_lag_step <- stepAIC(lm_model_month_lag,
                                  scope = list(upper = ~Time+Month_1+Month_2+
                                               Month_3+Month_4+Month_5+Month_6+Month_7+
                                               Month_8+Month_9+Month_10+Month_11+Month_12,
                                               lower = ~1), trace = TRUE,
                                  direction="backward")

colnames(lm_model_month_lag_step$model)
set.seed(12345)

data_oil <- astsa::oil
data_gas <- astsa::gas

ts.plot(data_oil, data_gas, gpars = list(col = c("black", "red")))

#alternative

autoplot(ts(cbind(data_oil, data_gas), start = 2000, frequency = 52)) +
        ylab("Price of Oil and Gas") +xlab("Year") +
        ggtitle("Price of Oil and Gas vs. Years")
set.seed(12345)

autoplot(ts(cbind(log(data_oil), log(data_gas)), start = 2000, frequency = 52)) +
        ylab("Log of price of Oil and Gas") +xlab("Year") +
        ggtitle("Log of price of Oil and Gas vs. Years")
set.seed(12345)

autoplot(ts(diff(data_oil, differences = 1), start = 2000, frequency = 52)) +
        ylab("Price of Oil") +xlab("Year") +
        ggtitle("Price of Oil with Diff 1 vs. Years")

autoplot(ts(diff(data_gas, differences = 1), start = 2000, frequency = 52)) +
        ylab("Price of Gas") +xlab("Year") +
        ggtitle("Price of Gas with Diff 1 vs. Years")
```

```r
ggAcf(diff(data_oil, differences = 1), data_oil)
ggAcf(diff(data_gas, differences = 1), data_gas)

set.seed(12345)

oil_price_one_diff <- diff(data_oil, differences = 1)
gas_price_one_diff <- diff(data_gas, differences = 1)

df <- data.frame(oil_price_one_diff=as.matrix(oil_price_one_diff),
            gas_price_one_diff = as.matrix(gas_price_one_diff),
                    time=time(oil_price_one_diff))

df <- na.omit(df)


df$gas_price_one_diff = lag(df$gas_price_one_diff,1)
df$gas_price_two_diff = lag(df$gas_price_one_diff,2)
df$gas_price_three_diff = lag(df$gas_price_one_diff,3)
df$oil_price_one_diff = lag(df$oil_price_one_diff,1)
df$oil_price_two_diff = lag(df$oil_price_one_diff,2)
df$oil_price_three_diff = lag(df$oil_price_one_diff,3)

df <- na.omit(df)

df$smooth_one_week_lag <- ksmooth(x = df$oil_price_one_diff, y = df$gas_price_one_diff, bandwidth = 0.4
df$smooth_two_week_lag <- ksmooth(x = df$oil_price_two_diff, y = df$gas_price_two_diff, bandwidth = 0.4
df$smooth_three_week_lag <- ksmooth(x = df$oil_price_three_diff, y = df$gas_price_three_diff, bandwidth


ggplot(data=df, aes(x=time)) +
  geom_line(aes(y= smooth_one_week_lag, color= "smooth_one_week_lag")) +
    geom_line(aes(y= smooth_two_week_lag, color= "smooth_two_week_lag")) +
    geom_line(aes(y= smooth_three_week_lag, color= "smooth_three_week_lag")) +
      scale_colour_manual("", breaks = c("smooth_one_week_lag", "smooth_two_week_lag", "smooth_three_we
                    values = c("#CC79A7", "#000000", "#D55E00")) +
  ggtitle("Smoothened Plot of lag")

set.seed(12345)

df$x_t_more_zero <- ifelse(df$oil_price_one_diff>0,"1","0")
lm_model_lag <- lm(data=df, formula = gas_price_one_diff~oil_price_one_diff+oil_price_two_diff)

plot(lm_model_lag$residuals, type = 'l')
ggAcf(lm_model_lag$residuals)
```